

Building large de Bruijn graphs



Giulio Ermanno Pibiri
Ca' Foscari University of Venice



Venice, Italy, 2 December 2025

Agenda


1. Context, motivations, and problem definition
2. A simple algorithm
3. Refined algorithms:
 - **BCALM** – minimizers and union-find
 - **GGCAT** – super-kmers and randomization

1. Context, motivations, and problem definition

Massive DNA Collections



- **Peta bytes of data available:**
 - ENA (European Nucleotide Archive)
 - SRA (Sequence Read Archive)
 - RefSeq (Reference Sequence Database)
 - Ensembl
- These collections are paving the way to answer fundamental questions regarding biology and evolution.



k-mers

- **Q.** But how do we exploit such potential?

We need efficient methods to index and search data at this scale.

- One popular strategy: transform a DNA sequence into a **set** of short substrings of fixed length k – the so-called **k-mers**.

ACGGTAGAACCGATTCAAATTGACGTAGC...

ACGGTAGAACCGA

CGGTAGAACCGAT

GGTAGAACCGATT

GTAGAACCGATTTC

TAGAACCGATTCA

AGAACCGATTCAA

GAACCGATTCAA

AACCGATTCAAAT

...

← Example for $k=13$.

k-mers applications

- Software tools based on k-mers are predominant in bioinformatics.
- Many applications:
 - genome assembly
 - variant calling
 - sequence comparison/alignment
 - pan-genome analysis
 - meta-genomics
 - ...
- But we will not talk about applications in the following...

Collapsing the redundancy in large k-mer sets

- Note that, given a set S of long strings, the corresponding k-mer set is **highly redundant** (for a suitable value of k).
- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.

Set of k-mers:

ACGT	CGTT
GTAA	TTAC
TACG	ACGA
CGAA	GAAA
CGAC	GACA
ACAA	CAAA
AAAT	AATT

Collapsing the redundancy in large k-mer sets

- Note that, given a set S of long strings, the corresponding k-mer set is **highly redundant** (for a suitable value of k).
- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.

Set of k-mers:

ACGT	CGTT
GTAA	TTAC
TACG	ACGA
CGAA	GAAA
CGAC	GACA
ACAA	CAAA
AAAT	AATT

- S alone would cost $12 + 11 + 10 + (2) = 35$ chars.
- But its 14 distinct k-mers cost $14 \cdot (4 + 1) = 70$ chars! Twice as much. This is due to **(k-1)-length overlaps** being represented redundantly.

A special char to
distinguish the strings

Collapsing the redundancy in large k-mer sets

- Note that, given a set S of long strings, the corresponding k-mer set is **highly redundant** (for a suitable value of k).
- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.

Set of k-mers:


ACGT	CGTT
GTAA	TTAC
TACG	ACGA
CGAA	GAAA
CGAC	GACA
ACAA	CAAA
AAAT	AATT

- S alone would cost $12 + 11 + 10 + (2) = 35$ chars.
- But its 14 distinct k-mers cost $14 \cdot (4 + 1) = 70$ chars! Twice as much. This is due to **(k-1)-length overlaps** being represented redundantly.
- **Q.** How to collapse (i.e., reduce) this redundancy?

A special char to distinguish the strings


de Bruijn graphs

- **de Bruijn graph.** A de Bruijn graph (dBG) of order $k > 0$ for a set S of strings is a directed graph $G_k(S)$ where nodes are the distinct k -mers of S and there exists a directed edge from x to y if $x[2..k] = y[1..k - 1]$ and x “glued” with y , i.e., $x + y[k]$, is a $(k+1)$ -mer of S .
- A path in the graph spells a string obtained by glueing all its k -mers. With a little abuse of notation we refer to paths and their spelled strings interchangeably.
- Example for $S = \{ \text{“ACGTTACGTTAC”}, \text{“ACGTTACGAAA”}, \text{“ACGACAAATT”} \}$ and $k = 4$.




de Bruijn graphs

- **de Bruijn graph.** A de Bruijn graph (dBG) of order $k > 0$ for a set S of strings is a directed graph $G_k(S)$ where nodes are the distinct k -mers of S and there exists a directed edge from x to y if $x[2..k] = y[1..k - 1]$ and x “glued” with y , i.e., $x + y[k]$, is a $(k+1)$ -mer of S .
- A path in the graph spells a string obtained by glueing all its k -mers. With a little abuse of notation we refer to paths and their spelled strings interchangeably.
- Example for $S = \{ \text{“ACGTTACGTTAC”}, \text{“ACGTTACGAAA”}, \text{“ACGACAAATT”} \}$ and $k = 4$.




de Bruijn graphs

- **de Bruijn graph.** A de Bruijn graph (dBG) of order $k > 0$ for a set S of strings is a directed graph $G_k(S)$ where nodes are the distinct k -mers of S and there exists a directed edge from x to y if $x[2..k] = y[1..k - 1]$ and x “glued” with y , i.e., $x + y[k]$, is a $(k+1)$ -mer of S .
- A path in the graph spells a string obtained by glueing all its k -mers. With a little abuse of notation we refer to paths and their spelled strings interchangeably.
- Example for $S = \{ \text{“ACGTTACGTTAC”}, \text{“ACGTTACGAAA”}, \text{“ACGACAAATT”} \}$ and $k = 4$.




de Bruijn graphs

- **de Bruijn graph.** A de Bruijn graph (dBG) of order $k > 0$ for a set S of strings is a directed graph $G_k(S)$ where nodes are the distinct k -mers of S and there exists a directed edge from x to y if $x[2..k] = y[1..k - 1]$ and x “glued” with y , i.e., $x + y[k]$, is a $(k+1)$ -mer of S .
- A path in the graph spells a string obtained by glueing all its k -mers. With a little abuse of notation we refer to paths and their spelled strings interchangeably.
- Example for $S = \{ \text{“ACGTTACGTTAC”}, \text{“ACGTTACGAAA”}, \text{“ACGACAAATT”} \}$ and $k = 4$.




Compacted de Bruijn graphs

- **Unitig and maximal unitig.** A unitig in $G_k(S)$ is a path where all inner nodes have in/out degree 1. A maximal unitig is a unitig that cannot be extended without loosing the property of being a unitig. Let U be the set of all maximal unitigs of $G_k(S)$.
- **Compacted de Bruijn graph.** The compacted dBG of $G_k(S)$ is a directed graph where nodes are the strings of U and there exists a directed edge from x to y if $x[|x| - k + 2..|x|] = y[1..k - 1]$ and $x[|x| - k + 1..|x|] + y[k]$ is a $(k+1)$ -mer in S .
- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.



Compacted de Bruijn graphs


- **Unitig and maximal unitig.** A unitig in $G_k(S)$ is a path where all inner nodes have in/out degree 1. A maximal unitig is a unitig that cannot be extended without loosing the property of being a unitig. Let U be the set of all maximal unitigs of $G_k(S)$.
- **Compacted de Bruijn graph.** The compacted dBG of $G_k(S)$ is a directed graph where nodes are the strings of U and there exists a directed edge from x to y if $x[|x| - k + 2..|x|] = y[1..k - 1]$ and $x[|x| - k + 1..|x|] + y[k]$ is a $(k+1)$ -mer in S .
- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.



$$U = \{ \text{"ACGTTACG"}, \text{"ACGA"}, \text{"CGAAA"}, \text{"CGACAAATT"} \}$$

Compacted de Bruijn graphs

- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.




$$U = \{ \text{"ACGTTACG"}, \text{"ACGA"}, \\ \text{"CGAAA"}, \text{"CGACAAATT"} \}$$

- Two remarks:**

Compacted de Bruijn graphs

- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.




$$U = \{ \text{"ACGTTACG"}, \text{"ACGA"}, \\ \text{"CGAAA"}, \text{"CGACAAATT"} \}$$

- Two remarks:**
 1. To store U we need $26+3=29$ chars, much better than the 70 chars for the “plain” k-mer set.

Compacted de Bruijn graphs

- Example for $S = \{ \text{"ACGTTACGTTAC"}, \text{"ACGTTACGAAA"}, \text{"ACGACAAATT"} \}$ and $k = 4$.



$$U = \{ \text{"ACGTTACG"}, \text{"ACGA"}, \text{"CGAAA"}, \text{"CGACAAATT"} \}$$

$$U^* = \{ \text{"ACGTTACGAAA"}, \text{"CGACAAATT"} \}$$

- Two remarks:
 - To store U we need $26+3=29$ chars, much better than the 70 chars for the “plain” k-mer set.
 - We could do even better by glueing some unitigs via a (smallest, i.e., with minimum number of paths) **disjoint-node path cover U^*** (just 20 chars!).

Is it relevant/useful?

“AllTheBacteria”, <https://allthebacteria.org>


- Example.

The **top-3** species in the “AllTheBacteria” collection take **1.3 TB of gzipped files**.

The maximal unitigs of the corresponding dBG take **< 9 GB**.

Hence, a $\approx 145 \times$ reduction.

- Save storage space and speed up applications involving large k-mer sets.



Taken from <https://www.biorxiv.org/content/10.1101/2024.03.08.584059v3.full>

Problem definition

- **Problem.** Given a set S of strings and an integer $k > 0$, we want to compute the set U of maximal unitigs of $G_k(S)$ quickly and using little computer memory.
- Usually S is very large (e.g., hundreds of thousands or millions of genomes) and $k \in \{31, \dots, 63\}$.
- An intensively studied problem. A lot of research effort is actively spent on it.

BCALM [Chikhi et al., 2015]

BCALM v2 [Chikhi, Limasset, and Medvedev, 2016]

TwoPaCo [Minkin, Pham, and Medvedev, 2016]

Cuttlefish [Khan and Patro, 2021]

Cuttlefish v2 [Khan et al., 2022]

GGCAT [Cracco and Tomescu, 2023] ← state of the art

2. A simple algorithm


A simple algorithm

- **Idea.** Build $G_k(S)$ and visit it.
- We start a new path from a randomly chosen k-mer $x \in S$, and extend it as much as possible forward and backward **as long as the path is a unitig** (i.e., inner nodes have one predecessor and one successor only).


The resulting path is a **maximal** unitig.

- During the visit, we keep track of visited nodes.
- We repeat this process until all nodes have been visited.


A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo




A simple algorithm – Demo



A simple algorithm – Demo




A simple algorithm – Demo



ACG–TTAC–G

C–GACA–AATT


A simple algorithm – Demo



ACG–TTAC–G

C–GACA–AATT

A simple algorithm – Demo




ACG–TTAC–G

C–GACA–AATT

ACGA

A simple algorithm – Demo




ACG–TTAC–G

C–GACA–AATT

ACGA

A simple algorithm – Demo




ACG–TTAC–G

C–GACA–AATT

ACGA

A simple algorithm – Demo




ACG–TTAC–G

C–GACA–AATT

ACGA

C–GAAA

A simple algorithm – Demo



ACG–TTAC–G

C–GACA–AATT

ACGA

C–GAAA

Q. How would you implement it?

Representing the graph

- We create a hash table G such that, for each distinct k-mer $x \in S$,
 $G[x] = (pred(x), succ(x))$

where

$$\begin{aligned} pred(x) &= \{c \mid c + x[1..k-1] \text{ is a k-mer of } S\} \\ succ(x) &= \{c \mid x[2..k] + c \text{ is a k-mer of } S\} \end{aligned}$$

- Let:
 - N = the total number of k-mers in S
 - n = the number of distinct k-mers in S (nodes of the dBG)
 - m = number of edges in the dBG
- G is built in $\Theta(N)$ time and takes

$$\begin{aligned} O(n + \sum_{x \in G} (|pred(x)| + |succ(x)|)) &= O(n + m) \text{ space, because} \\ \sum_{x \in G} (|pred(x)| + |succ(x)|) &= 2m. \end{aligned}$$

```
7  def dbg(S, k):
8      G = {}
9      for s in S:
10         n = len(s)
11         for i in range(n - k + 1):
12             x = s[i:i+k]
13             pred, succ = G.setdefault(x, (set(), set()))
14             if i > 0:
15                 pred.add(s[i-1])
16             if i + k < n:
17                 succ.add(s[i+k])
18
19     return G
```

Visiting the graph

```
59  def compact(G, k):
60      visited = set()
61      U = []
62      for x in G:
63          if x in visited:
64              continue
65          visited.add(x);
66          fwd = extend_fwd(x, k, G, visited)
67          bwd = extend_bwd(x, k, G, visited)
68          u = bwd + x + fwd
69          U.append(u)
70      return U
```

- $\Theta(n + m)$ time to visit the graph.

```
27  def extend_fwd(x, k, G, visited):
28      path = ''
29      while True:
30          _, succ = G[x]
31          if len(succ) != 1:
32              break
33          nuc = next(iter(succ))
34          nxt = x[1:] + nuc
35          pred, _ = G[nxt]
36          if len(pred) != 1 or nxt in visited:
37              break
38          visited.add(nxt)
39          path = path + nuc
40          x = nxt
41      return path

43  def extend_bwd(x, k, G, visited):
44      path = ''
45      while True:
46          pred, _ = G[x]
47          if len(pred) != 1:
48              break
49          nuc = next(iter(pred))
50          prv = nuc + x[:k-1]
51          _, succ = G[prv]
52          if len(succ) != 1 or prv in visited:
53              break
54          visited.add(prv)
55          path = nuc + path
56          x = prv
57      return path
```

Putting all together

```
83 k = 4
84 S = ["ACGTTACGTTAC", "ACGTTACGAAA", "ACGACAAATT"]
85
86 G = dbg(S, k)
87 U = compact(G, k)
88 print(U) # ['ACGTTACG', 'ACGA', 'CGAAA', 'CGACAAATT']
```

- Summary: complexity (in both time and space) is **linear** in the size of the graph.
- The problem here is the **space!**
- Let's call this algorithm SIMPLE in the following.

3. Refined algorithms

Bird's eye view


- **Divide and conquer.** Split the k-mer set into buckets and work on them independently in parallel using SIMPLE.
- Note: There are many engineering aspects (e.g., multithreading, compressed I/Os, disk pipelining) that we will not discuss but contribute significantly to keep the overall running time low.

Implementation matters a lot!

BCALM

Chikhi, Limasset, and Medvedev, 2016


- **Idea.** Distribute k-mers into buckets and process the buckets (sub-graphs) independently in parallel using SIMPLE. A bucket is a file on disk.
- Not all buckets are loaded into main memory at the same time to run SIMPLE, hence allowing to scale to very large datasets using a prescribed amount of RAM.
- The algorithm takes as input a set of kmers, hence a first preprocessing step of k-mer counting is performed.



Challenges


1. How to guarantee that two k-mers sharing an overlap are placed in the same bucket?

In principle, we could define one bucket per $(k-1)$ -mer but there would be too many for typical values of k .




Challenges

- 1. How to guarantee that two k-mers sharing an overlap are placed in the same bucket?**
In principle, we could define one bucket per $(k-1)$ -mer but there would be too many for typical values of k .
- 2. In general, a maximal unitig can span more than one bucket.**
Some unitigs computed in different buckets must be further glued into a maximal unitig.



Challenges

- 1. How to guarantee that two k-mers sharing an overlap are placed in the same bucket?**
In principle, we could define one bucket per $(k-1)$ -mer but there would be too many for typical values of k .
- 2. In general, a maximal unitig can span more than one bucket.**
Some unitigs computed in different buckets must be further glued into a maximal unitig.
- 3. Load-balancing: buckets should have approx. the same size.** (We will not talk about this.)



Minimizers

- **Minimizer.** Given a k -mer x and an order \mathcal{O} over all ℓ -mers, the minimizer of length $\ell \leq k$ of x is the (leftmost) smallest ℓ -mer of x according to \mathcal{O} .
- Example for $x = \text{TCGATAGAAC}$ ($k = 10$), $\ell = 4$, and using $\mathcal{O} = \text{lexicographic order}$.

TCGA
CGAT
GATA
ATAG
TAGA
→ **AGAA**
GAAC

Minimizers

- **Minimizer.** Given a k -mer x and an order \mathcal{O} over all ℓ -mers, the minimizer of length $\ell \leq k$ of x is the (leftmost) smallest ℓ -mer of x according to \mathcal{O} .
- Example for $x = \text{TCGATAGAAC}$ ($k = 10$), $\ell = 4$, and using $\mathcal{O} = \text{lexicographic order}$.


TCGA
CGAT
GATA
ATAG
TAGA
→ AGAA
GAAC

TCGATAGAACCGATTCAAATTCGAT...
TCGATAGAAC
CGATAGAACC
GATAGAACCG
ATAGAACCGA
TAGAACCGAT
AGAACCGATT
GAACCGATTCA
AACCGATTCA
...

- **Property.** Consecutive k -mers tend to share the same minimizer.



Left/right minimizers

- **Left/right minimizer.** For a k -mer x , let $lm(x)$ and $rm(x)$ be the left and right minimizer of x , defined as the minimizer of length $\ell \leq k - 1$ of the $(k-1)$ -length prefix/suffix of x , $x[1..k - 1]$ and $x[2..k]$ respectively.




Left/right minimizers

- **Left/right minimizer.** For a k -mer x , let $lm(x)$ and $rm(x)$ be the left and right minimizer of x , defined as the minimizer of length $\ell \leq k - 1$ of the $(k-1)$ -length prefix/suffix of x , $x[1..k - 1]$ and $x[2..k]$ respectively.
- Now, if x and y have a $(k-1)$ -length overlap, they have at least a minimizer in common. (The viceversa is not true in general.)
- Here, surely $rm(x) = lm(y)$.




Distribution via left/right minimizers

- BCALM places k-mer x in bucket $i \in \{1..4^\ell\}$ if $i = lm(x)$, for a small $\ell > 0$ (like $\ell = 8$). If $j = rm(x)$, then x **also** goes to bucket j .
- The intuition is that, since consecutive k-mers tend to share the same minimizer, a bucket is likely to hold k-mers that are part of the same unitig.
- The parameter ℓ controls the size of the buckets.




Distribution via left/right minimizers

- BCALM places k-mer x in bucket $i \in \{1..4^\ell\}$ if $i = lm(x)$, for a small $\ell > 0$ (like $\ell = 8$). If $j = rm(x)$, then x **also** goes to bucket j .
- The intuition is that, since consecutive k-mers tend to share the same minimizer, a bucket is likely to hold k-mers that are part of the same unitig.
- The parameter ℓ controls the size of the buckets.
- Bucket i also contains k-mers x for which $rm(x) = j$. Such k-mers might potentially induce a **branch** with the k-mers belonging to bucket j (e.g., the k-mer z).
- We therefore only compact k-mers x and y in bucket i if they share an overlap **and** $i = rm(x) = lm(y)$. We do **not** compact x and y in the example. The k-mers x for which $lm(x) \neq rm(x)$ **will be at the end of a unitig**.




Distribution and compaction – Example



- Example for $k = 4$ and $\ell = 2$.
- (Empty buckets omitted.)


Distribution and compaction – Example




- Example for $k = 4$ and $\ell = 2$.
- (Empty buckets omitted.)



Distribution and compaction – Example



- Example for $k = 4$ and $\ell = 2$.
- (Empty buckets omitted.)




Linking non-maximal unitigs: union-find

- k-mers x for which $lm(x) \neq rm(x)$ exist in two copies, in bucket $lm(x)$ and in bucket $rm(x)$. They appear at the end of some unitig. These k-mers must be deduplicated to obtain the final maximal unitigs.
- The goal is to partition the unitigs into groups, each group containing the ones that should be compacted. Clearly, each group can be compacted independently in parallel.
- To achieve the partitioning, BCALM uses a **union-find data structure**: operation union is performed between the first and last k-mer of a unitig.

Linking non-maximal unitigs: union-find


- k-mers x for which $lm(x) \neq rm(x)$ exist in two copies, in bucket $lm(x)$ and in bucket $rm(x)$. They appear at the end of some unitig. These k-mers must be deduplicated to obtain the final maximal unitigs.
- The goal is to partition the unitigs into groups, each group containing the ones that should be compacted. Clearly, each group can be compacted independently in parallel.
- To achieve the partitioning, BCALM uses a **union-find data structure**: operation union is performed between the first and last k-mer of a unitig.



AACT	ACCC	ACCCTC	CCTC	CTAA
CTAA	AACTC		ACTC	CTAC


Linking non-maximal unitigs: union-find

- k-mers x for which $lm(x) \neq rm(x)$ exist in two copies, in bucket $lm(x)$ and in bucket $rm(x)$. They appear at the end of some unitig. These k-mers must be deduplicated to obtain the final maximal unitigs.
- The goal is to partition the unitigs into groups, each group containing the ones that should be compacted. Clearly, each group can be compacted independently in parallel.
- To achieve the partitioning, BCALM uses a **union-find data structure**: operation union is performed between the first and last k-mer of a unitig.



Linking non-maximal unitigs: union-find


- k-mers x for which $lm(x) \neq rm(x)$ exist in two copies, in bucket $lm(x)$ and in bucket $rm(x)$. They appear at the end of some unitig. These k-mers must be deduplicated to obtain the final maximal unitigs.
- The goal is to partition the unitigs into groups, each group containing the ones that should be compacted. Clearly, each group can be compacted independently in parallel.
- To achieve the partitioning, BCALM uses a **union-find data structure**: operation union is performed between the first and last k-mer of a unitig.



AACT	ACCC	ACCCTC	CCTC	CTAA
CTAA	AACTC	ACTC	ACTC	CTAC


Linking non-maximal unitigs: union-find

- k-mers x for which $lm(x) \neq rm(x)$ exist in two copies, in bucket $lm(x)$ and in bucket $rm(x)$. They appear at the end of some unitig. These k-mers must be deduplicated to obtain the final maximal unitigs.
- The goal is to partition the unitigs into groups, each group containing the ones that should be compacted. Clearly, each group can be compacted independently in parallel.
- To achieve the partitioning, BCALM uses a **union-find data structure**: operation union is performed between the first and last k-mer of a unitig.



Linking non-maximal unitigs: union-find


- k-mers x for which $lm(x) \neq rm(x)$ exist in two copies, in bucket $lm(x)$ and in bucket $rm(x)$. They appear at the end of some unitig. These k-mers must be deduplicated to obtain the final maximal unitigs.
- The goal is to partition the unitigs into groups, each group containing the ones that should be compacted. Clearly, each group can be compacted independently in parallel.
- To achieve the partitioning, BCALM uses a **union-find data structure**: operation union is performed between the first and last k-mer of a unitig.



AACT	ACCC	ACCCCTC	CCTC	CTAA
CTAA	AACTC		ACTC	CTAC

Linking non-maximal unitigs: union-find

- k-mers x for which $lm(x) \neq rm(x)$ exist in two copies, in bucket $lm(x)$ and in bucket $rm(x)$. They appear at the end of some unitig. These k-mers must be deduplicated to obtain the final maximal unitigs.
- The goal is to partition the unitigs into groups, each group containing the ones that should be compacted. Clearly, each group can be compacted independently in parallel.
- To achieve the partitioning, BCALM uses a **union-find data structure**: operation union is performed between the first and last k-mer of a unitig.




Performance of union-find algorithms

Algorithm	Union	Find
<i>quick-find</i>	n	1
<i>quick-union</i>	tree height	tree height
<i>weighted quick-union</i>	$\log n$	$\log n$
<i>weighted quick-union with path compression</i>	almost 1 (<i>inverse Ackermann</i>)	almost 1 (<i>inverse Ackermann</i>)

- Performance for n items.
- All algorithms use an array (or two) of size n as their data structure.
- From Chapter 1.5 of “Algorithms”, 4-th Ed., Sedgewick and Wayne.


Union-find implementation detail

- A complicating matter for BCALM is that is not trivial to assign integer ids to k-mers, **using little space**, to implement union-find.
- A classic hash table could be used but it also stores the k-mers themselves.



Union-find implementation detail

- A complicating matter for BCALM is that is not trivial to assign integer ids to k-mers, **using little space**, to implement union-find.
- A classic hash table could be used but it also stores the k-mers themselves.
- But **minimal perfect hashing** can be used.
- That is, we build a function h that maps all the n distinct k-mers into the integers $[1..n]$ without collisions and without storing them at all!
- Such functions can be built very efficiently and with space usage between 1.8 – 3.0 bits per key (very close to optimal).



GGCAT

Cracco and Tomescu, 2023

- A refined and finely-engineered version of BCALM.
- **Idea 1.** Merge the k-mer counting step with unitig compaction. A problem in BCALM is that the input set of k-mers must be first materialised on disk.
- This first step takes a non-trivial amount of time (although k-mer counting algorithms are very efficient) and the final input set can take a large amount of disk space.
- **Idea 2.** Avoid union-find data structure but use a randomised approach.
- Many other important practical improvements, e.g., intermediate files are written to disk compressed, better multithreading, better hash tables for SIMPLE, etc.

Super-k-mers

- GGCAT avoids the k-mer counting step of BCALM by splitting the input strings of S into **super-(k-1)-mers** and inserting unique k-mers in the hash table used by SIMPLE.
- **Super-k-mer.** Given a string, a super-k-mer is a **maximal** sequence of consecutive k-mers having the same minimizer.
- Example for $k = 13$ and $\ell = 4$:


ACGGTAG**AACC**GATTCAAATTGATCGATTAAATTAGAGCGATAAC...
ACGGTAG**AACC**GA
CGGTAG**AACC**GAT
GGTAG**AACC**GATT
GTAG**AACC**GATTG
TAG**AACC**GATTCA
AG**AACC**GATTCAA
GACCGATTCAAA
AACCGATT**AAAT**
...

super-k-mer



- A super-k-mer is a space-efficient representation for its set of k-mers.
- Buckets in GGCAT are made of **super-(k-1)-mers** and not of individual k-mers which reduces space on disk (by a lot!).


Linking non-maximal unitigs



- We create a list of pairs (first/last k-mer, unitig-id), sorted by k-mer:


$L = [(a,4), (c,2), (d,5), (d,7), (e,6), (f,7) (x,4), (x,3), (y,3), (y,1), (z,1), (z,2), (z,5), (z,6)].$

Linking non-maximal unitigs




- We create a list of pairs (first/last k-mer, unitig-id), sorted by k-mer:
 $L = [(a,4), (c,2), (d,5), (d,7), (e,6), (f,7), (x,4), (x,3), (y,3), (y,1), (z,1), (z,2), (z,5), (z,6)].$
- And only unitig-id pairs that should be merged are retained: $L = [(5,7), (3,4), (1,3)]$.
- At the beginning all unitig ids are marked as **unsealed**.

Linking non-maximal unitigs: randomization




$$L = [(5,7), (3,4), (1,3)]$$


Linking non-maximal unitigs: randomization



$$L = [(5,7), (3,4), (1,3)]$$



Linking non-maximal unitigs: randomization




$L = [(5,7), (3,4), (1,3)]$

$L = [(5,7), (3,4), (\mathbf{1},3)]$




Linking non-maximal unitigs: randomization




$L = [(5,7), (3,4), (1,3)]$

$L = [(5,7), (3,4), (\mathbf{1},3)]$




Linking non-maximal unitigs: randomization




$L = [(5,7), (3,4), (1,3)]$

$L = [(5,7), (3,4), (\mathbf{1},3)]$

$L = [(3,\mathbf{4})]$




Linking non-maximal unitigs: randomization




$L = [(5,7), (3,4), (1,3)]$

$L = [(5,7), (3,4), (\mathbf{1},3)]$

$L = [(\mathbf{3},\mathbf{4})]$



Linking non-maximal unitigs: randomization




$L = [(5,7), (3,4), (1,3)]$


$L = [(5,7), (3,4), (\mathbf{1},3)]$

$L = [(\mathbf{3},\mathbf{4})]$

$L = []$



Linking non-maximal unitigs: randomization




$L = [(5,7), (3,4), (1,3)]$

$L = [(5,7), (3,4), (\mathbf{1},3)]$

$L = [(\mathbf{3},4)]$

$L = []$



- Since each element from a pair is selected **uniformly at random**, given two pairs that should be merged, they will end up in the same bucket, with probability **at least 1/4**.
- Thus, unitigs that have to be linked will end up in the same bucket after just four steps in expectation.

Performance

A portion of Table 1 from [Khan et al., 2022].

			BCALM 2	CUTTLEFISH 2
	Dataset	k	Thread-count	
Human	Human	27	8	04 h 23 min (6.7) 01 h 13 min (3.2)
		16	8	04 h 58 min (8.9) 56 min (3.3)
		55	8	04 h 01 min (7.4) 02 h 20 min (3.5)
		16	16	04 h 26 min (10.5) 02 h 02 min (3.7)
	Human RNA-seq	27	8	02 h 58 min (3.8) 30 min (2.9)
		16	16	02 h 46 min (3.9) 20 min (3.0)
	Gut microbiome	27	16	02 h 34 min (7.7) 26 min (3.5)
		55	16	03 h 02 min (12.5) 44 min (4.0)

A portion of Table 2
from [Cracco and Tomescu, 2023].

Data set	Server	k	Cuttlefish 2	GGCAT
Human reads	<i>Small</i>	27	1 h:15 min (3.95 GB) [209 GB]	1 h:16 min (4.54 GB) [220 GB]
		63	2 h:07 min (4.23 GB) [140 GB]	1 h:03 min (7.11 GB) [156 GB]
Gut microbiome reads	<i>Small</i>	27	0 h:30 min (3.35 GB) [78 GB]	0 h:22 min (6.09 GB) [78 GB]
		63	1 h:08 min (3.86 GB) [107 GB]	0 h:19 min (5.42 GB) [51 GB]
<i>Salmonella</i> genomes (309 K)	<i>Small</i>	119	1 h:04 min (3.13 GB) [97 GB]	0 h:12 min (5.33 GB) [32 GB]
		27	6 h:59 min (4.38 GB) [1515 GB]	3 h:38 min (3.46 GB) [378 GB]
		63	12 h:02 min (3.88 GB) [1145 GB]	3 h:31 min (3.96 GB) [274 GB]
		119	17 h:07 min (3.95 GB) [1088 GB]	3 h:39 min (4.12 GB) [279 GB]
		255	77 h:58 min (4.82 GB) [1056 GB]	3 h:44 min (4.33 GB) [325 GB]

Performance

BCALM is outperformed by CUTTLEFISH which is, in turn, outperformed by GGCAT!

A portion of Table 2 from [Cracco and Tomescu, 2023].

A portion of Table 1 from [Khan et al., 2022].

Dataset	<i>k</i>	Thread-count	BCALM 2	CUTTLEFISH 2
Human	27	8	04 h 23 min (6.7)	01 h 13 min (3.2)
		16	04 h 58 min (8.9)	56 min (3.3)
	55	8	04 h 01 min (7.4)	02 h 20 min (3.5)
		16	04 h 26 min (10.5)	02 h 02 min (3.7)
Human RNA-seq	27	8	02 h 58 min (3.8)	30 min (2.9)
		16	02 h 46 min (3.9)	20 min (3.0)
Gut microbiome	27	16	02 h 34 min (7.7)	26 min (3.5)
	55		03 h 02 min (12.5)	44 min (4.0)

Data set	Server	<i>k</i>	Cuttlefish 2	GGCAT
Human reads	<i>Small</i>	27	1 h:15 min (3.95 GB) [209 GB]	1 h:16 min (4.54 GB) [220 GB]
		63	2 h:07 min (4.23 GB) [140 GB]	1 h:03 min (7.11 GB) [156 GB]
Gut microbiome reads	<i>Small</i>	27	0 h:30 min (3.35 GB) [78 GB]	0 h:22 min (6.09 GB) [78 GB]
		63	1 h:08 min (3.86 GB) [107 GB]	0 h:19 min (5.42 GB) [51 GB]
<i>Salmonella</i> genomes (309 K)	<i>Small</i>	119	1 h:04 min (3.13 GB) [97 GB]	0 h:12 min (5.33 GB) [32 GB]
		27	6 h:59 min (4.38 GB) [1515 GB]	3 h:38 min (3.46 GB) [378 GB]
		63	12 h:02 min (3.88 GB) [1145 GB]	3 h:31 min (3.96 GB) [274 GB]
		119	17 h:07 min (3.95 GB) [1088 GB]	3 h:39 min (4.12 GB) [279 GB]
		255	77 h:58 min (4.82 GB) [1056 GB]	3 h:44 min (4.33 GB) [325 GB]

References

1. Chikhi, Rayan, et al. "On the representation of de Bruijn graphs." *Journal of Computational Biology* 22.5 (2015): 336-352.
2. Chikhi, Rayan, Antoine Limasset, and Paul Medvedev. "Compacting de Bruijn graphs from sequencing data quickly and in low memory." *Bioinformatics* 32.12 (2016): i201-i208.
<https://github.com/GATB/bcalm>
3. Minkin, Ilia, Son Pham, and Paul Medvedev. "TwoPaCo: an efficient algorithm to build the compacted de Bruijn graph from many complete genomes." *Bioinformatics* 33.24 (2017): 4024-4032.
<https://github.com/medvedevgroup/TwoPaCo>
4. Khan J, Patro R. 2021. Cuttlefish: fast, parallel and low-memory compaction of de Bruijn graphs from large-scale genome collections. *Bioinformatics* 37: i177–i186. doi:10.1093/bioinformatics/btab309
<https://github.com/COMBINE-lab/cuttlefish>
5. Khan J, Kokot M, Deorowicz S, Patro R. 2022. Scalable, ultra-fast, and low- memory construction of compacted de Bruijn graphs with Cuttlefish 2. *Genome Biol* 23: 190. doi:10.1186/s13059-022-02743-6
<https://github.com/COMBINE-lab/cuttlefish>
6. Cracco, Andrea, and Alexandru I. Tomescu. "Extremely fast construction and querying of compacted and colored de Bruijn graphs with GGCAT." *Genome Research* 33.7 (2023): 1198-1207.
<https://github.com/algbio/GGCAT>