



Repetition-aware compression and query of colored de Bruijn graphs

Giulio Ermanno Pibiri

Ca' Foscari University of Venice



@giulio_pibiri



@jermp

RIKEN-AIP Workshop: Compressed Data Structures for Advanced Data Analysis
Tokyo, Japan, 4 July 2025

The colored k-mer indexing problem

- We are given a collection $\mathcal{R} = \{R_1, \dots, R_N\}$ of reference sequences. Each R_c is a (long) sequence over the DNA alphabet {A,C,G,T}.
- **Problem.** We want to build an *index* for \mathcal{R} so that we can retrieve the set $\text{ColorSet}(x) = \{c \mid x \in R_c\}$ efficiently for any k-mer x . Note that $\text{ColorSet}(x) = \emptyset$ if $x \notin \mathcal{R}$.

The colored k-mer indexing problem

- We are given a collection $\mathcal{R} = \{R_1, \dots, R_N\}$ of reference sequences. Each R_c is a (long) sequence over the DNA alphabet {A,C,G,T}.
- **Problem.** We want to build an *index* for \mathcal{R} so that we can retrieve the set $\text{ColorSet}(x) = \{c \mid x \in R_c\}$ efficiently for any k-mer x . Note that $\text{ColorSet}(x) = \emptyset$ if $x \notin \mathcal{R}$.
- A lot of hype in the indexing community for the case where \mathcal{R} is a **pangenome**, i.e., a collection of related genomes.
- **Applications.** This problem is relevant for applications where sequences are first matched against known references (i.e., mapping/alignment algorithms): single-cell RNA-seq, metagenomics, etc.


Modular indexing layout

- **Goal.** What we want is the map $x \rightarrow \text{ColorSet}(x) = \{c \mid x \in R_c\}$.
- **Two** data structures:

Modular indexing layout


- **Goal.** What we want is the map $x \rightarrow \text{ColorSet}(x) = \{c \mid x \in R_c\}$.
- **Two** data structures:
 1. All the distinct k-mers in $\mathcal{R} = \{R_1, \dots, R_N\}$ are stored in the **dictionary** \mathcal{D} .

\mathcal{D} stores n distinct k-mers and supports a $\text{Lookup}(x)$ operation which, given a k-mer x , returns a unique integer $1 \leq h \leq n$ if $x \in \mathcal{D}$, and \perp otherwise.



Modular indexing layout


- **Goal.** What we want is the map $x \rightarrow \text{ColorSet}(x) = \{c \mid x \in R_c\}$.
- **Two** data structures:
 1. All the distinct k-mers in $\mathcal{R} = \{R_1, \dots, R_N\}$ are stored in the **dictionary** \mathcal{D} .
 \mathcal{D} stores n distinct k-mers and supports a $\text{Lookup}(x)$ operation which, given a k-mer x , returns a unique integer $1 \leq h \leq n$ if $x \in \mathcal{D}$, and \perp otherwise.
 2. The sets $\{\text{ColorSet}(x)\}_x$ are stored in order of $\text{Lookup}(x)$ in the **inverted index** \mathcal{L} .



Modular indexing layout


- Many k-mer based indexes (all of them?) are incarnations/adaptations of this **modular indexing layout**, $\mathcal{D} + \mathcal{L}$:

- deBGA [Liu et al. 2016]
- Kallisto [Bray et al. 2016]
- BIGSI [Bradley et al. 2017]
- Rainbowfish [Almodaresi et al. 2017]
- Mantis [Pandey et al. 2018]
- Pufferfish [Almodaresi et al. 2018]
- SeqOthello [Yu et al. 2018]
- COBS [Bingmann et al. 2019]
- Reindeer [Marchet et al. 2020]
- Raptor [Seiler et al. 2021]
- Metagraph [Karasikov et al. 2022]
- NIQKI [Agret et al. 2022]
- Pufferfish2 [Fan et al. 2022]
- Themisto [Alanko et. al 2023]
- **Fulgor** [Fan et. al 2023, 2024; P. et al. 2024; Campanelli et al., 2024, 2025]




Modular indexing layout

- Our problem reduces to that of **representing two data structures**, \mathcal{D} and \mathcal{L} .
- To do so at best, we **must understand/exploit the properties** of our problem.
- **Q.** What are these properties?




de Bruijn graphs

- The dictionary \mathcal{D} is a set of k-mers with (k-1)-symbol overlaps.
- One-to-one correspondence between \mathcal{D} and a dBG.
- Example for **k = 3**.




Colored de Bruijn graphs

- Example for $k = 3$ and $N = 6$ references. References in \mathcal{R} are spelled by **paths** in the graph.




Colored de Bruijn graphs

- Example for $k = 3$ and $N = 6$ references. References in \mathcal{R} are spelled by **paths** in the graph.




Colored de Bruijn graphs

- Example for $k = 3$ and $N = 6$ references. References in \mathcal{R} are spelled by **paths** in the graph.



Colored compacted de Bruijn graphs

- Example for $k = 3$ and $N = 6$ references.
- Nodes having the **same color set along non-branching paths** are collapsed into **monochromatic unitigs**.



- Let's now index this object!

Properties of colored compacted dBGs

1. Unitigs spell references in \mathcal{R} .

- We can **represent the set of unitigs** instead of the set of k-mers: represent \mathcal{D} with **SSHash**. Better space effectiveness and cache locality (seen on Wed 2 July).


\mathcal{D}	
u_1	CAT
u_2	GACA
u_3	CGAACG
u_4	CGTCCG
u_5	AGCG
u_6	ATTAT
u_7	GAGTT
u_8	CGGAT
u_9	ATAGA
u_{10}	CGCTCG

Properties of colored compacted dBGs

2. Unitigs are monochromatic.

- We store the color set of each unitig, rather than for each k-mer because $\text{ColorSet}(x) = \text{ColorSet}(y)$ if k-mers x and y are part of the same unitig.
Thus, we need an efficient **map from k-mers to unitigs**: $x \rightarrow \text{Unitig}(x)$.

- For SSHash, it is easy to compute the **unitig identifier** $\text{Unitig}(x)$ given the k-mer x .
- Now \mathcal{L} stores $\text{ColorSet}(x)$ for each unitig in the order given by $\text{Unitig}(x)$.




Properties of colored compacted dBGs

2. Unitigs are monochromatic.

- We store the color set of each unitig, rather than for each k-mer because $\text{ColorSet}(x) = \text{ColorSet}(y)$ if k-mers x and y are part of the same unitig.
Thus, we need an efficient **map from k-mers to unitigs**: $x \rightarrow \text{Unitig}(x)$.

- For SSHash, it is easy to compute the **unitig identifier** $\text{Unitig}(x)$ given the k-mer x .
- Now \mathcal{L} stores $\text{ColorSet}(x)$ for each unitig in the order given by $\text{Unitig}(x)$.



Properties of colored compacted dBGs

3. **Unitigs co-occur.** → Distinct unitigs often have the same color set, i.e., they co-occur in the same subset of references. We have **way less distinct color sets than unitigs**.
We need an efficient **map from unitigs to distinct color sets**: $\text{Unitig}(x) \rightarrow \text{ColorSet}(x)$.

\mathcal{D}	
u_1	CAT
u_2	GACA
u_3	CGAACG
u_4	CGTCCG
u_5	AGCG
u_6	ATTAT
u_7	GAGTT
u_8	CGGAT
u_9	ATAGA
u_{10}	CGCTCG

Properties of colored compacted dBGs

3. **Unitigs co-occur.** → Distinct unitigs often have the same color set, i.e., they co-occur in the same subset of references. We have **way less distinct color sets than unitigs**.
We need an efficient **map from unitigs to distinct color sets**: $\text{Unitig}(x) \rightarrow \text{ColorSet}(x)$.


- Remember from Wed 2 July: SSHash stores a set of unitigs **in any wanted order (order-preserving)**.
- We can thus **permute** the unitigs in \mathcal{D} so that **consecutive unitigs have the same color set**.
- Then, mapping a unitig to its color is as simple as a **rank query** over a bitmap (see next).

\mathcal{D}	
u_1	CAT
u_2	GACA
u_3	CGAACG
u_4	CGTCCG
u_5	AGCG
u_6	ATTAT
u_7	GAGTT
u_8	CGGAT
u_9	ATAGA
u_{10}	CGCTCG


Properties of colored compacted dBGs

3. **Unitigs co-occur.** → Distinct unitigs often have the same color set, i.e., they co-occur in the same subset of references. We have **way less distinct color sets than unitigs**.
We need an efficient **map from unitigs to distinct color sets**: $\text{Unitig}(x) \rightarrow \text{ColorSet}(x)$.

- Remember from Wed 2 July: SSHash stores a set of unitigs **in any wanted order (order-preserving)**.
- We can thus **permute** the unitigs in \mathcal{D} so that **consecutive unitigs have the same color set**.
- Then, mapping a unitig to its color is as simple as a **rank query** over a bitmap (see next).



Mapping unitigs to color sets in succinct space



Compression of color sets

- Any method to compress integer sequences would work.

Techniques for inverted index compression

P. and Venturini, 2020, ACM Computing Surveys

- Here, we used three simple encodings based on the density of a set (ratio $|C|/N$):
 - If $|C|/N < 1/4$, the set is **sparse** and we compress the gaps with Elias' δ .
 - If $|C|/N \geq 3/4$, the set is **very dense** and we compress the gaps of \bar{C} (complementary set) with Elias' δ .
 - Otherwise, the set is **dense** and we code it with a bitmap of N bits.
 - (The thresholds $1/4$ and $3/4$ are optimal.)

Results

⚡ Fulgor (v1) <https://github.com/jermp/fulgor/releases/tag/v1.0.0> (2023, 2024)

Genomes	Rate	Fulgor		Themisto		MetaG.-B		MetaG.-NB		COBS		
		mm:ss	GB	h:mm:ss	GB	mm:ss	GB	h:mm:ss	GB	h:mm:ss	GB	
EC	3,682	98.99	2:10	1.68	0:03:40	2.46	22:00	30.44	1:05:41	0.40	0:45:11	34.93
	5,000	89.49	1:16	0.82	0:03:50	1.82	14:14	36.54	0:20:32	0.33	0:38:34	41.93
	10,000	89.71	2:26	2.11	0:07:35	4.16	28:15	92.18	0:43:40	0.61	1:01:14	84.20
SE	50,000	91.25	19:15	18.53	0:42:02	33.14	—	—	4:30:03	2.72	3:54:18	408.82
	100,000	91.41	27:30	42.78	1:22:00	75.93	—	—	9:40:06	4.82	8:07:29	522.56
	150,000	91.52	42:30	70.55	2:00:13	124.27	—	—	—	—	7:47:14	522.63
GB	30,691	92.91	01:10	30.02	0:01:20	48.47	28:55	15.86	0:22:05	9.91	0:34:45	225.57

Full-intersection query mode; 16 processing threads; Max RSS in GB; output of /usr/bin/time.

Yet another property!

1. Unitigs spell references in \mathcal{R} .
2. Unitigs are monochromatic.
3. Unitigs co-occur.
4. **Color sets are similar when indexing pangenomes.** → Opportunity to achieve **much better compression** if color sets are not compressed individually (each set independently of the others) but **common patterns are factored out and compressed once**.

Color sets are similar when indexing pangenomes

$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$

$C_2 = [2, 3, 15]$

$C_3 = [1, 3, 5, 7, 9, 10, 11]$

$C_4 = [1, 3, 5, 7, 9, 11, 13]$

$C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$

$C_6 = [6, 8]$

$C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$

$C_8 = [12, 16]$

- The pattern $\{3, 5, 9, 11\}$ is currently represented **three times**.
- The pattern $\{1, 11, 12, 13, 14, 16\}$ is represented **twice**.

Color sets are similar when indexing pangenesomes

$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$

$C_2 = [2, 3, 15]$

$C_3 = [1, 3, 5, 7, 9, 10, 11]$

$C_4 = [1, 3, 5, 7, 9, 11, 13]$

$C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$

$C_6 = [6, 8]$

$C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$

$C_8 = [12, 16]$

- The pattern $\{3, 5, 9, 11\}$ is currently represented **three times**.
- The pattern $\{1, 11, 12, 13, 14, 16\}$ is represented **twice**.
- C_3 and C_4 are very similar.

Color sets are similar when indexing pangenesomes

$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$

$C_2 = [2, 3, 15]$

$C_3 = [1, 3, 5, 7, 9, 10, 11]$

$C_4 = [1, 3, 5, 7, 9, 11, 13]$

$C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$

$C_6 = [6, 8]$

$C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$

$C_8 = [12, 16]$

- The pattern $\{3, 5, 9, 11\}$ is currently represented **three times**.
- The pattern $\{1, 11, 12, 13, 14, 16\}$ is represented **twice**.
- C_3 and C_4 are very similar.
- **Q.** How to **factor out** this redundancy?

Two ways of partitioning the color sets

$$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$$

$$C_2 = [2, 3, 15]$$

$$C_3 = [1, 3, 5, 7, 9, 10, 11]$$

$$C_4 = [1, 3, 5, 7, 9, 11, 13]$$

$$C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$$

$$C_6 = [6, 8]$$

$$C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$$

$$C_8 = [12, 16]$$

(a) Three horizontal partitions

$$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$$

$$C_2 = [2, 3, 15]$$

$$C_3 = [1, 3, 5, 7, 9, 10, 11]$$

$$C_4 = [1, 3, 5, 7, 9, 11, 13]$$

$$C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$$

$$C_6 = [6, 8]$$


$$C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$$

$$C_8 = [12, 16]$$

(b) Four vertical partitions

Horizontal partitioning: diff. and repr. color sets

- Example for **N = 16** references and **4** partitions.



Vertical partitioning: meta and partial color sets

- Example for **N = 16** references and **4 partitions**.

$\{ 1 12 13 14 16 \} \{ 3 5 9 \} \{ 7 11 \} \{ 2 4 6 8 10 15 \}$
new identifiers $\rightarrow 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16$

this defines a
permutation π

\mathcal{L}

$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$
 $C_2 = [2, 3, 15]$
 $C_3 = [1, 3, 5, 7, 9, 10, 11]$
 $C_4 = [1, 3, 5, 7, 9, 11, 13]$
 $C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$
 $C_6 = [6, 8]$
 $C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$
 $C_8 = [12, 16]$

π

$C_1 = [3|6, 7, 8|10|12, 15, 16]$
 $C_2 = [6|11, 16]$
 $C_3 = [1|6, 7, 8|9, 10|15]$
 $C_4 = [1, 3|6, 7, 8|9, 10]$
 $C_5 = [1, 2, 3, 4, 5|6, 8|9, 10|13]$
 $C_6 = [13, 14]$
 $C_7 = [1, 2, 3, 4, 5|6|10|14]$
 $C_8 = [2, 5]$

Vertical partitioning: meta and partial color sets

- Example for **N = 16** references and **4 partitions**.

new identifiers → 1 2 3 4 5 { 1 12 13 14 16 } { 3 5 9 } { 7 11 } { 2 4 6 8 10 15 } ← this defines a
partition 2 permutation π

\mathcal{L}

$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$
 $C_2 = [2, 3, 15]$
 $C_3 = [1, 3, 5, 7, 9, 10, 11]$
 $C_4 = [1, 3, 5, 7, 9, 11, 13]$
 $C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$
 $C_6 = [6, 8]$
 $C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$
 $C_8 = [12, 16]$

π

$C_1 = [3|6, 7, 8|10|12, 15, 16]$
 $C_2 = [6|11, 16]$
 $C_3 = [1|6, 7, 8|9, 10|15]$
 $C_4 = [1, 3|6, 7, 8|9, 10]$
 $C_5 = [1, 2, 3, 4, 5|6, 8|9, 10|13]$
 $C_6 = [13, 14]$
 $C_7 = [1, 2, 3, 4, 5|6|10|14]$
 $C_8 = [2, 5]$

Vertical partitioning: meta and partial color sets

- Example for **N = 16** references and **4 partitions**.

$\{ 1 \ 12 \ 13 \ 14 \ 16 \} \{ 3 \ 5 \ 9 \} \{ 7 \ 11 \} \{ 2 \ 4 \ 6 \ 8 \ 10 \ 15 \}$

new identifiers → 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16

partition 2

← this defines a permutation π

\mathcal{L}


```
 $C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$ 
 $C_2 = [2, 3, 15]$ 
 $C_3 = [1, 3, 5, 7, 9, 10, 11]$ 
 $C_4 = [1, 3, 5, 7, 9, 11, 13]$ 
 $C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$ 
 $C_6 = [6, 8]$ 
 $C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$ 
 $C_8 = [12, 16]$ 
```

π

```
 $C_1 = [3|6, 7, 8|10|12, 15, 16]$ 
 $C_2 = [6|11, 16]$ 
 $C_3 = [1|6, 7, 8|9, 10|15]$ 
 $C_4 = [1, 3|6, 7, 8|9, 10]$ 
 $C_5 = [1, 2, 3, 4, 5|6, 8|9, 10|13]$ 
 $C_6 = [13, 14]$ 
 $C_7 = [1, 2, 3, 4, 5|6|10|14]$ 
 $C_8 = [2, 5]$ 
```

Vertical partitioning: meta and partial color sets

- Example for **N = 16** references and **4 partitions**.



\mathcal{L}

```
 $C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$ 
 $C_2 = [2, 3, 15]$ 
 $C_3 = [1, 3, 5, 7, 9, 10, 11]$ 
 $C_4 = [1, 3, 5, 7, 9, 11, 13]$ 
 $C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$ 
 $C_6 = [6, 8]$ 
 $C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$ 
 $C_8 = [12, 16]$ 
```

π

$C_1 = [3|6, 7, 8|10|12, 15, 16]$
 $C_2 = [6|11, 16]$
 $C_3 = [1|6, 7, 8|9, 10|15]$
 $C_4 = [1, 3|6, 7, 8|9, 10]$
 $C_5 = [1, 2, 3, 4, 5|6, 8|9, 10|13]$
 $C_6 = [13, 14]$
 $C_7 = [1, 2, 3, 4, 5|6|10|14]$
 $C_8 = [2, 5]$

[6,7,8]
[6]
[6,8]

Vertical partitioning: meta and partial color sets

- Example for **N = 16** references and **4 partitions**.

$\{ 1 \ 12 \ 13 \ 14 \ 16 \} \{ 3 \ 5 \ 9 \} \{ 7 \ 11 \} \{ 2 \ 4 \ 6 \ 8 \ 10 \ 15 \}$

new identifiers → 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16

partition 2

← this defines a permutation π

\mathcal{L}

$C_1 = [3, 4, 5, 9, 10, 11, 13, 15]$
 $C_2 = [2, 3, 15]$
 $C_3 = [1, 3, 5, 7, 9, 10, 11]$
 $C_4 = [1, 3, 5, 7, 9, 11, 13]$
 $C_5 = [1, 3, 6, 7, 9, 11, 12, 13, 14, 16]$
 $C_6 = [6, 8]$
 $C_7 = [1, 3, 8, 11, 12, 13, 14, 16]$
 $C_8 = [12, 16]$


π

$C_1 = [3|6, 7, 8|10|12, 15, 16]$
 $C_2 = [6|11, 16]$
 $C_3 = [1|6, 7, 8|9, 10|15]$
 $C_4 = [1, 3|6, 7, 8|9, 10]$
 $C_5 = [1, 2, 3, 4, 5|6, 8|9, 10|13]$
 $C_6 = [13, 14]$
 $C_7 = [1, 2, 3, 4, 5|6|10|14]$
 $C_8 = [2, 5]$

[6,7,8]
[6]
[6,8]
↓ -5
[1,2,3]
[1]
[1,3]


Vertical partitioning: meta and partial color sets

- Example for $N = 16$ references and 4 partitions.




Vertical partitioning: meta and partial color sets

- Example for $N = 16$ references and 4 partitions.



Vertical partitioning: meta and partial color sets

- Example for $N = 16$ references and 4 partitions.



Results – Index size


 **Fulgor** (v3) <https://github.com/jermp/fulgor/releases/tag/v3.0.0> (2024)

Dataset	Fulgor			d-Fulgor		m-Fulgor		md-Fulgor	
	dBG	Color sets	Total	Color sets	Total	Color sets	Total	Color sets	Total
EC	0.29	1.36 (83%)	1.65	0.45 (61%)	0.74	0.40 (58%)	0.69	0.24 (45%)	0.52
SE-5K	0.16	0.59 (79%)	0.75	0.20 (56%)	0.36	0.16 (50%)	0.32	0.11 (40%)	0.27
SE-10K	0.35	1.66 (83%)	2.01	0.48 (58%)	0.83	0.34 (49%)	0.70	0.22 (39%)	0.57
SE-50K	1.25	17.03 (93%)	18.29	4.31 (77%)	5.57	2.08 (62%)	3.34	1.38 (52%)	2.64
SE-100K	1.71	40.71 (96%)	42.43	9.37 (84%)	11.10	3.75 (68%)	5.47	2.26 (57%)	3.98
SE-150K	2.02	68.61 (97%)	70.65	15.73 (89%)	17.77	5.27 (72%)	7.31	3.22 (61%)	5.26
GB	21.29	15.54 (42%)	36.83	7.51 (26%)	28.81	9.16 (30%)	30.46	6.19 (23%)	27.48

Save money!

Amazon EC2 instances pricing:

- <https://instances.vantage.sh/aws/ec2/x2gd.medium>
16 GiB of RAM – **73 \$** per month
- <https://instances.vantage.sh/aws/ec2/x2gd.xlarge>
64 GiB of RAM – **292 \$** per month
- <https://instances.vantage.sh/aws/ec2/x2gd.2xlarge>
128 GiB of RAM – **584 \$** per month
- <https://instances.vantage.sh/aws/ec2/x2gd.4xlarge>
256 GiB of RAM – **1168 \$** per month



md-Fulgor: **5.3 GB**

m-Fulgor: **7.3 GB**

d-Fulgor: **18 GB**

Fulgor: **71 GB**

Original data: **225 GB**

(150,000 S. Enterica genomes,
compressed with gzip)

Numbers taken on 25/06/2025.

Results – Query time

⚡Fulgor (v3) <https://github.com/jermp/fulgor/releases/tag/v3.0.0> (2024)

Dataset	Mapping rate	h-Fulgor		d-Fulgor		m-Fulgor		md-Fulgor	
		GB	GB	GB	GB	GB	GB	GB	GB
EC	98.99	2:12	1.67	4:52	0.78	3:08	0.73	6:07	0.57
SE-5K	89.49	1:14	0.80	1:54	0.41	1:25	0.37	2:10	0.32
SE-10K	89.71	2:29	2.06	4:14	0.90	2:56	0.77	4:55	0.65
SE-50K	91.25	14:05	18.24	27:25	5.82	17:00	3.64	33:25	2.95
SE-100K	91.41	29:00	42.40	58:10	11.58	34:40	6.08	1:09:00	4.63
SE-150K	91.52	44:30	70.55	1:31:00	18.55	53:00	8.29	1:50:00	6.29
GB	92.91	1:10	36.01	1:00	28.25	1:09	29.79	1:03	26.88

Full-intersection query mode; 16 processing threads; Max RSS in GB; output of /usr/bin/time.

Faster query times!


⚡Fulgor (v4) <https://github.com/jermp/fulgor/releases/tag/v4.0.0> (2025)

Dataset	Mapping rate	h-Fulgor			d-Fulgor			m-Fulgor			md-Fulgor		
		before	after	GB	before	after	GB	before	after	GB	before	after	GB
EC	98.99	2:12	2:10	1.67	4:52	2:29	0.78	3:08	1:32	0.73	6:07	1:41	0.57
SE-5K	89.49	1:14	1:10	0.80	1:54	1:44	0.41	1:25	1:09	0.37	2:10	1:21	0.32
SE-10K	89.71	2:29	2:20	2.06	4:14	2:54	0.90	2:56	2:07	0.77	4:55	2:30	0.65
SE-50K	91.25	14:05	12:00	18.24	27:25	14:50	5.82	17:00	10:10	3.64	33:25	11:50	2.95
SE-100K	91.41	29:00	24:00	42.40	58:10	29:50	11.58	34:40	20:30	6.08	1:09:00	22:50	4.63
SE-150K	91.52	44:30	37:00	70.55	1:31:00	44:20	18.55	53:00	30:20	8.29	1:50:00	33:50	6.29
GB	92.91	1:10	1:10	36.01	1:00	1:26	28.25	1:09	1:02	29.79	1:03	1:02	26.88


Full-intersection query mode; 16 processing threads; Max RSS in GB; output of /usr/bin/time.

Faster query times!


Fulgor (v4) <https://github.com/jermp/fulgor/releases/tag/v4.0.0> (2025)



(a) SE-50K



(b) SE-100K



(c) SE-150K


Full-intersection query mode; 16 processing threads; Max RSS in GB; output of /usr/bin/time.

Faster query times – Intuition (for intersection)


- If most color sets being intersected are **super-dense**: compute the union of complement sets, say U , and iterate over the complement of U .
- If color sets are relative to the **same representative color set**, their intersection can be computed **only using their symmetric set differences** (that are much shorter).
- Meta and partial color sets yield a **two-level** index organisation:
we can first intersect meta color sets and only compute intersections of partial color sets for the common partitions only.
- Other tricks for (threshold-based) union of color sets apply.
- Check our new paper <https://jermp.github.io/assets/pdf/papers/WABI2025.pdf> out!

Repetition-aware improvement factors

- Results obtained by indexing *Salmonella Enterica* genomes from the “AllTheBacteria” collection [Hunt et al., 2024], <https://allthebacteria.org>



(a) Space improvement



(b) Query time improvement

Full-intersection query mode; 16 processing threads

Conclusions

- SSHash to obtain an **efficient map from k-mers to unitigs** (2022).
- Permute unitigs in color set order to enable a **space-efficient mapping from unitigs to color sets** (2023).
- Color sets can be factorized into “patterns” that **capture their repetitiveness** (2024): two paradigms – **horizontal and vertical partitioning** – that can be combined.
- These patterns can also be exploited to achieve faster query times (2025).

Conclusions

- SSHash to obtain an **efficient map from k-mers to unitigs** (2022).
- Permute unitigs in color set order to enable a **space-efficient mapping from unitigs to color sets** (2023).
- Color sets can be factorized into “patterns” that **capture their repetitiveness** (2024): two paradigms – **horizontal and vertical partitioning** – that can be combined.
- These patterns can also be exploited to achieve faster query times (2025).

A special thank to my co-authors!



Rob Patro



Alessio Campanelli

Next goals

- Index more and more and make the indexes available to the community, .e.g.
- the entire “AllTheBacteria”
<https://allthebacteria.org>
- Logan, <https://github.com/IndexThePlanet/Logan>
- All NCBI viruses
<https://www.ncbi.nlm.nih.gov/labs/virus/vssi>
- etc.
- **Sharding to external memory** seems inevitable in the future.

Taken from <https://www.biorxiv.org/content/10.1101/2024.03.08.584059v3.full>

