


# U-index: A Universal Indexing Framework for Matching Long Patterns

**Giulio Ermanno Pibiri**

Ca' Foscari University of Venice



@jermmp.bsky.social



@jermmp

**23-rd Symposium on Experimental Algorithms (SEA 2025)**

Venice, Italy, 22 July 2025

Joint work with

**Lorraine A. K. Ayad**

**Gabriele Fici**

**Ragnar Groot Koerkamp**

**Grigorios Loukides**

**Rob Patro**

**Solon P. Pissis**

# The text indexing problem

- Given a string  $T[0..n)$  over the alphabet  $\Sigma$ , pre-process  $T$  so that the following queries can be answered efficiently for any string  $P[0..m)$ :
  - *Locate*( $P, T$ ): return all the positions where  $P$  occurs in  $T$ ;
  - *Count*( $P, T$ ): count the number of occurrences of  $P$  in  $T$ ;
  - *Extract*( $i, j, T$ ): report the substring  $T[i..j]$ .
- Fundamental and well-studied problem.

# The text indexing problem

- Given a string  $T[0..n)$  over the alphabet  $\Sigma$ , pre-process  $T$  so that the following queries can be answered efficiently for any string  $P[0..m)$ :
  - *Locate*( $P, T$ ): return all the positions where  $P$  occurs in  $T$ ;
  - *Count*( $P, T$ ): count the number of occurrences of  $P$  in  $T$ ;
  - *Extract*( $i, j, T$ ): report the substring  $T[i..j]$ .
- Fundamental and well-studied problem.
- **More about this on Friday: “25 years of compressed self-indexes” !**

# The text indexing problem

- Many solutions with different trade-offs between space and time.
- Solutions broadly fall into two categories:

# The text indexing problem

- Many solutions with different trade-offs between space and time.
- Solutions broadly fall into two categories:
  1. **Compressed:** The text is replaced (is “self-indexed”) with a compressed representation.
  2. **Uncompressed:** A redundancy (an “index”) is attached to  $T$  to accelerate queries.

# The text indexing problem

- Many solutions with different trade-offs between space and time.
- Solutions broadly fall into two categories:
  - 1. Compressed:** The text is replaced (is “self-indexed”) with a compressed representation.
  - 2. Uncompressed:** A redundancy (an “index”) is attached to  $T$  to accelerate queries.
- Solutions in 1. are very space-efficient but generally slower to build and query than solutions in 2. which — on the other hand — are space-inefficient.
- **Example.** The (uncompressed) **suffix array** is much faster to query than the FM-index but requires  $n \log n$  bits on top of the text.

# Our contribution

- We focus on the uncompressed case, i.e., we attach an index to the text, and address the space-inefficiency of the index while supporting efficient queries.

# Our contribution

- We focus on the uncompressed case, i.e., we attach an index to the text, and address the space-inefficiency of the index while supporting efficient queries.
- **Main idea:** if we compute a **sketch** of the text  $T$ , say  $S = \text{Sketch}(T)$ , then  $\text{Index}(S)$  will be smaller/faster than  $\text{Index}(T)$  because  $S$  is **a lot smaller** than  $T$ , for any  $\text{Index}$ .
- At query time: we also compute  $Q = \text{Sketch}(P)$  and match  $Q$  against  $S$ . Candidate matches (including *false positives*) are mapped back to  $T$  to be verified.



# Our contribution

- We focus on the uncompressed case, i.e., we attach an index to the text, and address the space-inefficiency of the index while supporting efficient queries.
- **Main idea:** if we compute a **sketch** of the text  $T$ , say  $S = \text{Sketch}(T)$ , then  $\text{Index}(S)$  will be smaller/faster than  $\text{Index}(T)$  because  $S$  is **a lot smaller** than  $T$ , for any  $\text{Index}$ .
- At query time: we also compute  $Q = \text{Sketch}(P)$  and match  $Q$  against  $S$ . Candidate matches (including *false positives*) are mapped back to  $T$  to be verified.
- That is, we have a **universal framework** because:
  - **any index** can be used for  $S$ ;
  - **any locally-consistent sampling algorithm** can be used to sketch the text and obtain  $S$ .

# Intermezzo: sketching with minimizers

- Consider each window of  $w$  consecutive  $k$ -mers from a string  $T$ : sample one  $k$ -mer out of  $w$  and call it the “representative” of the window — or its *minimizer*.
- We would like to sample the **same minimizer** from consecutive windows so that the **set of distinct minimizers** forms a succinct sketch for  $T$ .

Example for  $w = 4$  and  $k = 7$ .

ACGGTAGAACCGATTCAAATTCGAT...

ACGGTAGAAC  
CGGTAGAAC  
GGTAGAACCG  
GTAGAACCGA  
TAGAACCGAT  
AGAACCGATT  
GAACCGATTTC  
AACCGATTCA  
...

# Intermezzo: sketching with minimizers

- **Q.** How do we compare different sampling algorithms?

**A.** We define the *density* of a sampling algorithm as the fraction between the number of (distinct) minimizers and the total number of  $k$ -mers of  $T$ .

**The lower the density, the better!**

# Intermezzo: sketching with minimizers

- **Q.** How do we compare different sampling algorithms?

**A.** We define the *density* of a sampling algorithm as the fraction between the number of (distinct) minimizers and the total number of  $k$ -mers of  $T$ .

**The lower the density, the better!**

- Since the same  $k$ -mer cannot be a minimizer for more than  $w$  consecutive  $k$ -mers, we immediately have a **lower bound** of  $1/w$  on the density of any sampling algorithm.

```
TAGAACCGAT
AGAACCGATT
GAACCGATTTC
AACCGATTCA
...
```

# The “folklore”, random, minimizer

- Take the **leftmost smallest  $k$ -mer** of the window according to an order  $\mathcal{O}_k$ .
- We usually define the total order using a random hash function (*random* minimizer).
- In this case, the density is  $2/(w + 1)$ : almost a factor of 2 away from the lower bound for large  $w$ .

---

```
1: function MINIMIZER( $W, w, k, \mathcal{O}_k$ )
2:    $o_{min} = +\infty$ 
3:    $p = 0$ 
4:   for  $i = 0; i < w; i = i + 1$  do
5:      $o = \mathcal{O}_k(W[i..i + k])$ 
6:     if  $o < o_{min}$  then
7:        $o_{min} = o$ 
8:        $p = i$ 
9:   return  $p$ 
```

---

# The “folklore”, random, minimizer

- Take the **leftmost smallest  $k$ -mer** of the window according to an order  $\mathcal{O}_k$ .
- We usually define the total order using a random hash function (*random* minimizer).
- In this case, the density is  $2/(w + 1)$ : almost a factor of 2 away from the lower bound for large  $w$ .

---

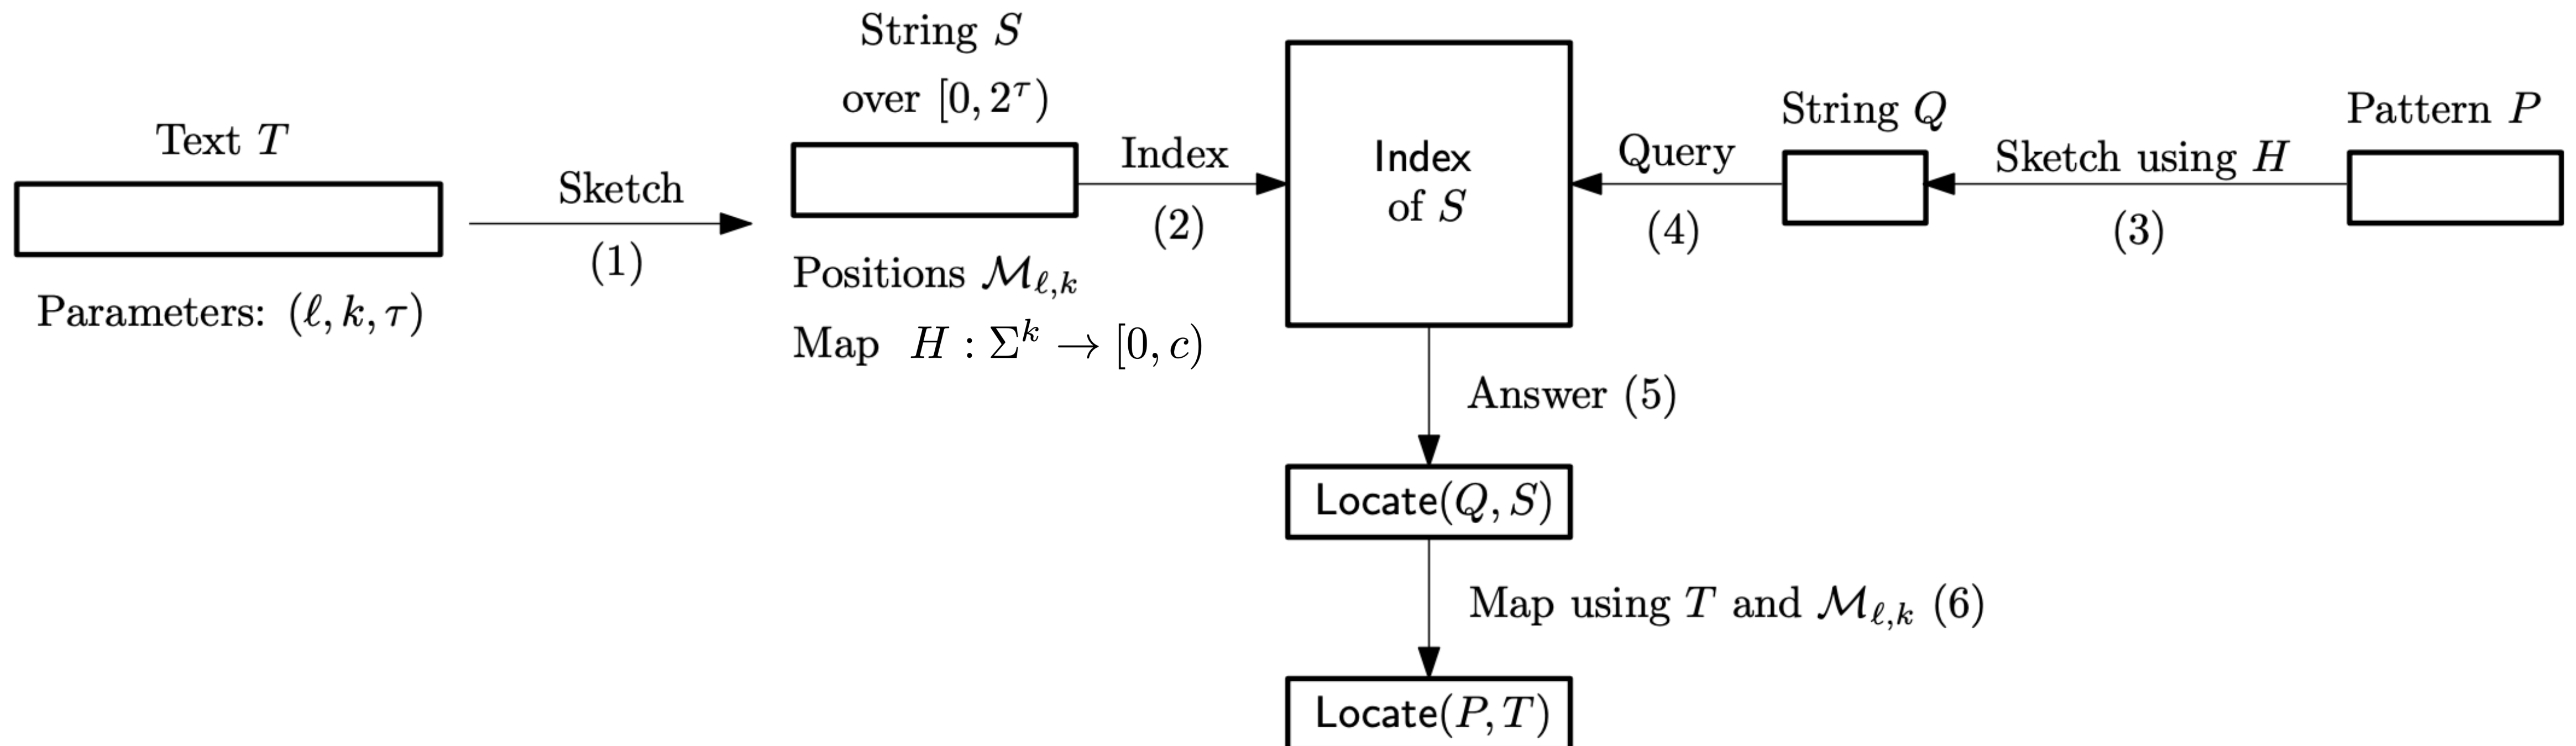
```
1: function MINIMIZER( $W, w, k, \mathcal{O}_k$ )
2:    $o_{min} = +\infty$ 
3:    $p = 0$ 
4:   for  $i = 0; i < w; i = i + 1$  do
5:      $o = \mathcal{O}_k(W[i..i + k])$ 
6:     if  $o < o_{min}$  then
7:        $o_{min} = o$ 
8:        $p = i$ 
9:   return  $p$ 
```

---

**More about random minimizers on Thursday!**

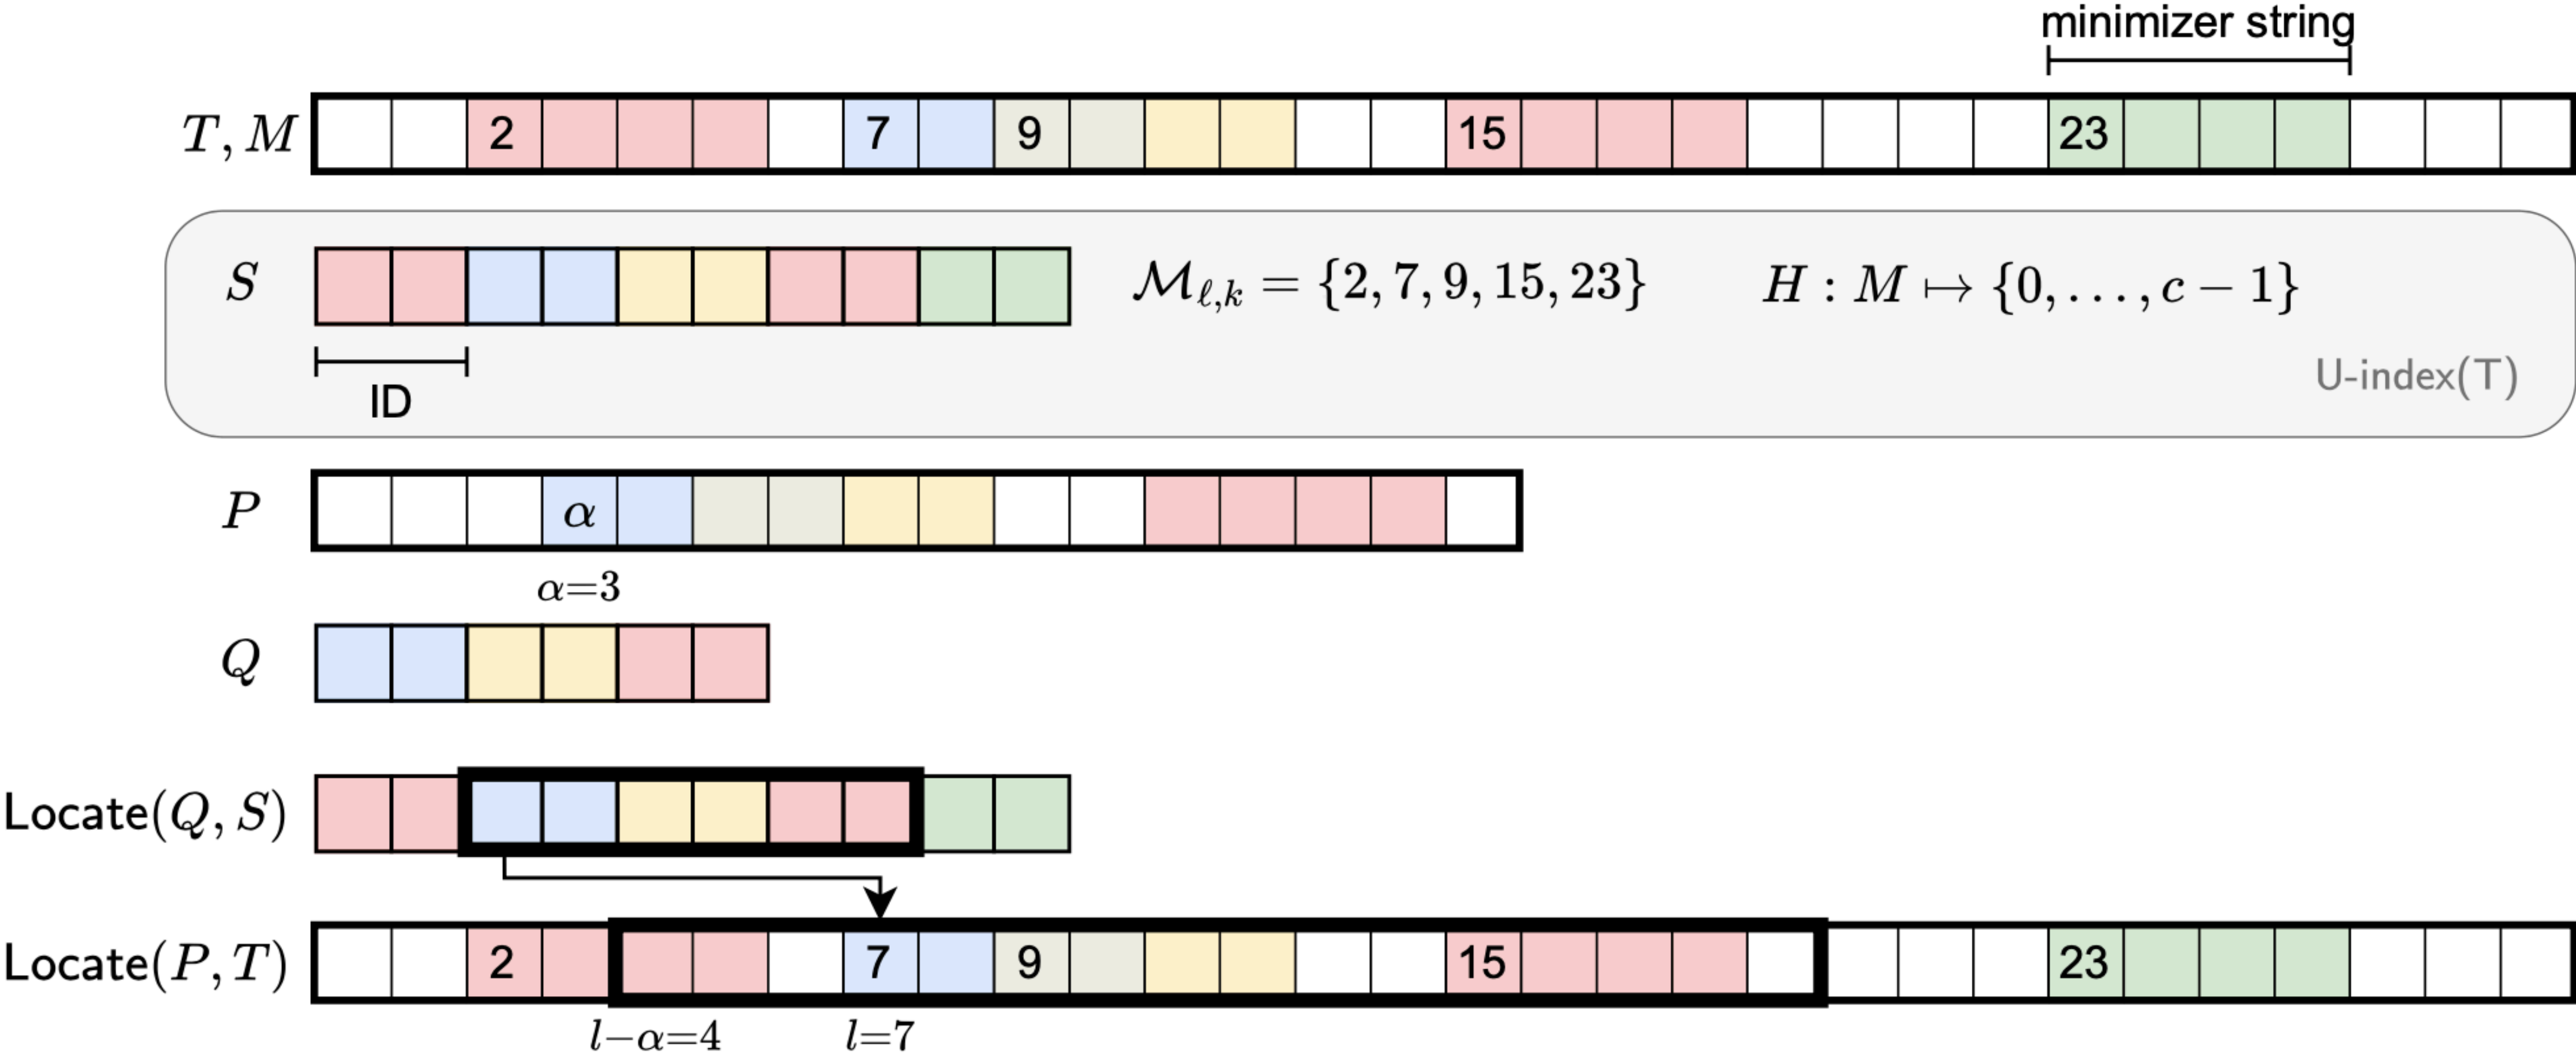
# The U-index framework for matching long patterns

- We fix integers  $k > 0$  and  $\ell \geq k$  and let  $w := \ell - k + 1$ , so that any pattern  $P$  of length  $m \geq \ell$  contains at least one minimizer.






# An example






# Theoretical guarantees

- Using some machinery, we guarantee that an occurrence in  $Locate(Q, S)$  is verified in  $O(1)$ , rather than  $O(m)$ . This can be done in  $O(z)$  space on top of the space of the text, where  $z$  is the number of minimizers of  $T$ .




# Theoretical guarantees

- Using some machinery, we guarantee that an occurrence in  $Locate(Q, S)$  is verified in  $O(1)$ , rather than  $O(m)$ . This can be done in  $O(z)$  space on top of the space of the text, where  $z$  is the number of minimizers of  $T$ .




# Theoretical guarantees

- Using some machinery, we guarantee that an occurrence in  $Locate(Q, S)$  is verified in  $O(1)$ , rather than  $O(m)$ . This can be done in  $O(z)$  space on top of the space of the text, where  $z$  is the number of minimizers of  $T$ .




# Theoretical guarantees

- Using some machinery, we guarantee that an occurrence in  $Locate(Q, S)$  is verified in  $O(1)$ , rather than  $O(m)$ . This can be done in  $O(z)$  space on top of the space of the text, where  $z$  is the number of minimizers of  $T$ .




# Theoretical guarantees


- Using some machinery, we guarantee that an occurrence in  $Locate(Q, S)$  is verified in  $O(1)$ , rather than  $O(m)$ . This can be done in  $O(z)$  space on top of the space of the text, where  $z$  is the number of minimizers of  $T$ .



# Results — Index size and build space for human chr 1



# Results — Build and query time for human chr 1



# Conclusions

- Main take-away: U-index is a framework to **enhance the performance of any off-the-shelf text index**, provided that the patterns to match are **reasonably long**.
- Example application: **long read mapping** to reference genomes.
- **Bottleneck: verifying false positive matches.**
- Rust code: <https://github.com/u-index/u-index-rs>



# Conclusions

- Main take-away: U-index is a framework to **enhance the performance of any off-the-shelf text index**, provided that the patterns to match are **reasonably long**.
- Example application: **long read mapping** to reference genomes.
- **Bottleneck: verifying false positive matches.**
- Rust code: <https://github.com/u-index/u-index-rs>


**Thank you for the attention!**  
**A special thank to all my co-authors!**

# An example application


- We demonstrated that the U-index framework can be useful for **long read mapping**.
- A core problem in Computational Biology; it involves aligning long patterns to a reference genome.
- Experimental setting: we align 450 HiFi long reads (avg. length is 16 kbp) on a complete human reference genome. We use  $k = 8$  and  $\ell = 128$  and split each read in patterns of length  $m = 256$ .

# An example application

- We demonstrated that the U-index framework can be useful for **long read mapping**.
- A core problem in Computational Biology; it involves aligning long patterns to a reference genome.
- Experimental setting: we align 450 HiFi long reads (avg. length is 16 kbp) on a complete human reference genome. We use  $k = 8$  and  $\ell = 128$  and split each read in patterns of length  $m = 256$ .
- Very practical numbers **using a suffix array** as index: the U-index is built in 12 seconds with  $\approx 9\mu\text{s}$  per pattern (23 avg. false positives per pattern).




# Density by varying $k$




- Example for  $w = 24$ .
- Measured over a string of 10 million i.i.d. random characters with an alphabet size of 4.
- <https://github.com/jermp/minimizers>

# Density by varying $k$



- Example for  $w = 24$ .
- Measured over a string of 10 million i.i.d. random characters with an alphabet size of 4.
- <https://github.com/jermp/minimizers>

# Density by varying $k$



- Example for  $w = 24$ .
- Measured over a string of 10 million i.i.d. random characters with an alphabet size of 4.
- <https://github.com/jermp/minimizers>